

Ecosystem Service Change & Hotspot Extraction

vv0.5.0

Jerónimo Rodríguez Escobar

2025-12-16 14:58 Pacific Standard Time

Table of contents

1 Overview	1
2 Setup and Libraries	2
3 Metadata Banner	2
4 Add regional attributes to Grid (countries & biomes)	2
5 Produce signed change bars	3
5.1 Diagnose sign flips (keep0 vs drop0)	6
6 Load cached long table (plt_long) for hotspots/KS	8
7 Hotspot extraction workflow (global + subregional)	8
7.1 Hotspot rules & export configuration	9
7.2 Validate hotspot configuration	10
7.3 Export hotspot layers	10
8 Trimmed change bar plots	10
8.1 TODO: Hotspot vs non-hotspot contribution bars	11
8.2 Signed means (alt variants: keep0 vs drop0; dashed = global mean)	11
9 Hotspot violin plots	21

1 Overview

Prepares the processed 10 km grid, derives hotspot flags, and produces the bar/violin plots that summarize signed and absolute change by group (region, income, biome). Key steps: load helpers and paths, attach grouping attributes, pivot change fields to long form, define hotspots using the configured thresholds, export hotspot layers, and write the latest plots to outputs/plots/latest/.

2 Setup and Libraries

Load all core packages (tidyverse, spatial, plotting helpers) and initialize reproducibility settings. This chunk also sources `R(paths.R)` and runs `devtools::load_all()` so downstream chunks can reuse helper functions.

3 Metadata Banner

Published: 2025-12-17 13:10 PST



Run metadata

- Analysis version: v0.5.0
- Rendered: 2025-12-17 13:10 PST
- Git: feat/ks-analysis @ 10146f9
- Data root: /home/jeronimo/data/global_ncp

4 Add regional attributes to Grid (countries & biomes)

Attach country and biome IDs to the 10 km grid (`sf_f`) so downstream aggregation/hotspot steps can group by region. This chunk only runs if the processed GPKG is missing.

The input 10 km grid (`sf_f`) is enriched with country and WWF biome IDs to allow aggregation and comparisons of change by subregions (World Bank region, income group, continent, UN region, and biome). We use a point-on-surface join to avoid sliver/overlap issues, keep only the needed fields, and write a single enriched GPKG for downstream grouping, hotspot extraction, and plotting.



What this step does

- Reads country and biome layers from `vectors/`.
- Joins them to the 10 km grid via point-on-surface.
- Writes `processed/10k_change_calc.gpkg` for downstream analysis (pivoting, hotspots, plots).
Inputs: `sf_f` grid; `vectors/cartographic_ee_ee_r264_correspondence.gpkg`; `vectors/Biome.gpkg`
Output: `processed/10k_change_calc.gpkg` (grid + regional attributes)

Why point-on-surface? It's robust for odd cell shapes and avoids polygon–polygon sliver issues.

Skipping join: /home/jeronimo/data/global_ncp/processed/10k_change_calc.gpkg already exists.

5 Produce signed change bars

Bars are computed directly from `processed/10k_change_calc.gpkg`. We output two variants—zeros kept vs zeros dropped—both with trimmed tails and a dashed global mean. Pivoted `plt_long` is only necessary for hotspots/KS.

What this chunk does:

- Read the processed GPKG and normalize IDs (`fid`, `c_fid`), dropping any `c_fid.x/y` artifacts.
- Build a service mapping (`svc_map`) so raw column bases (with/without `_mean`) map to the canonical services; construct `col_pct`/`col_abs` names.
- `agg_by_group()`: intersect requested groupings with available columns; for each grouping \times service, pick the pct/abs change column, handle Inf pct values (`handle_inf`), optionally drop zeros (`drop_zero`), trim extremes at `cut_q`, then compute mean change by group. Returns tidy rows with `grouping`, `group`, `service`, `metric`, `mean_chg`.
- Run `agg_by_group()` four ways: trimmed keep0/drop0 and no-trim keep0/drop0 \rightarrow `signed_keep0`, `signed_drop0`, `signed_keep0_no_trim`, `signed_drop0_no_trim`.
- A later chunk computes global reference means; plotting is in its own chunk so you can skip it when debugging.

```
services <- svc_order
groupings <- c("income_grp", "region_wb", "continent", "region_un", "WWF_biome")
cut_q <- 0.999
handle_inf <- "na" # options: "na", "cap"
verbose <- TRUE
# helper for lightweight progress messages

vmsg <- function(...) if (isTRUE(verbose)) message(...)

gpkg <- file.path(Sys.getenv("GLOBAL_NCP_DATA"), "processed", "10k_change_calc.gpkg")
stopifnot(file.exists(gpkg))

vmsg("Reading: ", pkg, " [layer: 10k_change_calc]")
sf_f <- sf::st_read(gpkg, layer = "10k_change_calc", quiet = TRUE)
sf_f <- sf::st_drop_geometry(sf_f)
if (!"fid" %in% names(sf_f)) sf_f$fid <- seq_len(nrow(sf_f))
if (!"c_fid" %in% names(sf_f)) {
  if ("c_fid.x" %in% names(sf_f)) sf_f <- dplyr::rename(sf_f, c_fid = c_fid.x)
  else if ("c_fid.y" %in% names(sf_f)) sf_f <- dplyr::rename(sf_f, c_fid = c_fid.y)
  else if ("id" %in% names(sf_f)) sf_f <- dplyr::rename(sf_f, c_fid = id)
}
sf_f <- dplyr::select(sf_f, -dplyr::any_of(c("c_fid.x", "c_fid.y")))

# build mapping; include both raw base and no_mean base, with ready-to-use columns
chg_cols <- grep("_abs|pct)_chg$", names(sf_f), value = TRUE)
vmsg("Found ", length(chg_cols), " change columns; mapping services...")
services_raw <- unique(sub("_abs|pct)_chg$", "", chg_cols)) # as in file (may include _mean)
services_clean <- stringr::str_remove(services_raw, "_mean$")
services_lower <- tolower(services_clean)
svc_map <- tibble::tibble(
  col_base = c(services_raw, services_clean),
```

```

canonical = dplyr::recode(c(services_lower, services_lower),
                           !!!canonical_lookup, .default = c(services_clean, services_clean))
) |>
  dplyr::distinct() |>
  dplyr::mutate(col_pct = paste0(col_base, "_pct_chg"),
                col_abs = paste0(col_base, "_abs_chg"))

agg_by_group <- function(df, services, groupings, cut_q = 0.999,
                           handle_inf = c("na", "cap"), drop_zero = FALSE) {
  handle_inf <- match.arg(handle_inf)
  groupings <- intersect(groupings, names(df))
  vmsg("Groupings available: ", paste(groupings, collapse = ", "))
  out <- list()
  for (g in groupings) {
    vmsg(" Grouping: ", g)
    for (svc in services) {
      map_rows <- dplyr::filter(svc_map, canonical == svc)
      if (!nrow(map_rows)) next
      cols_pct <- map_rows$col_pct[map_rows$col_pct %in% names(df)]
      cols_abs <- map_rows$col_abs[map_rows$col_abs %in% names(df)]
      if (!length(cols_pct) && !length(cols_abs)) next

      add_one <- function(col, metric) {
        v <- df[[col]]
        if (metric == "pct" && any(is.infinite(v))) {
          if (handle_inf == "na") v[is.infinite(v)] <- NA_real_
        }
        if (isTRUE(drop_zero)) v[v == 0] <- NA_real_
        cap <- stats::quantile(abs(v), cut_q, na.rm = TRUE)
        v_trim <- pmax(pmin(v, cap), -cap)
        tibble::tibble(
          service = svc,
          group = df[[g]],
          metric = metric,
          val = v_trim
        )
      }

      rows <- list()
      if (length(cols_pct)) rows[[length(rows)+1]] <- add_one(cols_pct[1], "pct")
      if (length(cols_abs)) rows[[length(rows)+1]] <- add_one(cols_abs[1], "abs")
      rows <- dplyr::bind_rows(rows)
      if (!nrow(rows)) next
      rows <- rows |>
        dplyr::filter(!is.na(group)) |>
        dplyr::group_by(service, metric, group) |>
        dplyr::summarise(mean_chg = mean(val, na.rm = TRUE), .groups = "drop") |>

```

```

        dplyr::mutate(grouping = g)
        out[[length(out)+1]] <- rows
    }
}
dplyr::bind_rows(out)
}

# trimmed means, zeros kept/dropped
signed_keep0 <- agg_by_group(sf_f, services, groupings, cut_q = cut_q,
                               handle_inf = handle_inf, drop_zero = FALSE)
signed_drop0 <- agg_by_group(sf_f, services, groupings, cut_q = cut_q,
                               handle_inf = handle_inf, drop_zero = TRUE)
# untrimmed (cut_q = 1) means, zeros kept/dropped
signed_keep0_no_trim <- agg_by_group(sf_f, services, groupings, cut_q = 1,
                                       handle_inf = handle_inf, drop_zero = FALSE)
signed_drop0_no_trim <- agg_by_group(sf_f, services, groupings, cut_q = 1,
                                       handle_inf = handle_inf, drop_zero = TRUE)
print(dplyr::count(signed_keep0, grouping, metric, service))

```

```

# A tibble: 80 x 4
  grouping metric service      n
  <chr>    <chr>  <chr>     <int>
1 WWF_biome abs   C_Risk       16
2 WWF_biome abs   C_Risk_Red_Ratio 16
3 WWF_biome abs   N_Ret_Ratio   16
4 WWF_biome abs   N_export     16
5 WWF_biome abs   Nature_Access 16
6 WWF_biome abs   Pollination   16
7 WWF_biome abs   Sed_Ret_Ratio 16
8 WWF_biome abs   Sed_export    16
9 WWF_biome pct   C_Risk       16
10 WWF_biome pct  C_Risk_Red_Ratio 16
# i 70 more rows

```

```
print(dplyr::count(signed_drop0, grouping, metric, service))
```

```

# A tibble: 80 x 4
  grouping metric service      n
  <chr>    <chr>  <chr>     <int>
1 WWF_biome abs   C_Risk       16
2 WWF_biome abs   C_Risk_Red_Ratio 16
3 WWF_biome abs   N_Ret_Ratio   16
4 WWF_biome abs   N_export     16
5 WWF_biome abs   Nature_Access 16
6 WWF_biome abs   Pollination   16
7 WWF_biome abs   Sed_Ret_Ratio 16

```

```

8 WWF_biome abs      Sed_export          16
9 WWF_biome pct      C_Risk              16
10 WWF_biome pct     C_Risk_Red_Ratio   16
# i 70 more rows

```

5.1 Diagnose sign flips (keep0 vs drop0)

Use this helper to see where dropping zeros changes the sign of the mean. Run interactively (set `eval: true`) after the direct bars chunk so `signed_keep0`/`signed_drop0` exist.

```

# Join trimmed keep/drop and untrimmed keep/drop, with clear column names
cmp <- signed_keep0 |>
  dplyr::rename(mean_keep_trim = mean_chg) |>
  dplyr::full_join(
    signed_drop0 |> dplyr::rename(mean_drop_trim = mean_chg),
    by = c("grouping", "metric", "service", "group")
  ) |>
  dplyr::full_join(
    signed_keep0_no_trim |> dplyr::rename(mean_keep_notrim = mean_chg),
    by = c("grouping", "metric", "service", "group")
  ) |>
  dplyr::full_join(
    signed_drop0_no_trim |> dplyr::rename(mean_drop_notrim = mean_chg),
    by = c("grouping", "metric", "service", "group")
  ) |>
  dplyr::mutate(
    sign_keep_trim = sign(mean_keep_trim),
    sign_drop_trim = sign(mean_drop_trim),
    sign_keep_notrim = sign(mean_keep_notrim),
    sign_drop_notrim = sign(mean_drop_notrim),
    flipped_trim = sign_keep_trim != sign_drop_trim &
      !is.na(sign_keep_trim) & !is.na(sign_drop_trim),
    flipped_notrim = sign_keep_notrim != sign_drop_notrim &
      !is.na(sign_keep_notrim) & !is.na(sign_drop_notrim)
  )

flips_trim_only <- cmp |> dplyr::filter(flipped_trim == TRUE)
flips_no_trim_only <- cmp |> dplyr::filter(flipped_notrim==TRUE)

```

```

# Trimmed, zeros kept: compare pct vs abs signs
flips_pct_abs_keep_trim <- signed_keep0 |>
  tidyr::pivot_wider(names_from = metric, values_from = mean_chg) |>
  dplyr::mutate(
    sign_pct = sign(pct),
    sign_abs = sign(abs),
    flipped_pct_abs = sign_pct != sign_abs &
      !is.na(sign_pct) & !is.na(sign_abs)
  )

```

```

) |>
dplyr::filter(flipped_pct_abs)

# Same for zeros dropped
flips_pct_abs_drop_trim <- signed_drop0 |>
  tidyr::pivot_wider(names_from = metric, values_from = mean_chg) |>
  dplyr::mutate(
    sign_pct = sign(pct),
    sign_abs = sign(abs),
    flipped_pct_abs = sign_pct != sign_abs &
      !is.na(sign_pct) & !is.na(sign_abs)
  ) |>
  dplyr::filter(flipped_pct_abs)

# Trimmed, zeros kept: compare pct vs abs signs
flips_pct_abs_keep_no_trim <- signed_keep0_no_trim |>
  tidyr::pivot_wider(names_from = metric, values_from = mean_chg) |>
  dplyr::mutate(
    sign_pct = sign(pct),
    sign_abs = sign(abs),
    flipped_pct_abs = sign_pct != sign_abs &
      !is.na(sign_pct) & !is.na(sign_abs)
  ) |>
  dplyr::filter(flipped_pct_abs)

# Same for zeros dropped
flips_pct_abs_drop_no_trim <- signed_drop0_no_trim |>
  tidyr::pivot_wider(names_from = metric, values_from = mean_chg) |>
  dplyr::mutate(
    sign_pct = sign(pct),
    sign_abs = sign(abs),
    flipped_pct_abs = sign_pct != sign_abs &
      !is.na(sign_pct) & !is.na(sign_abs)
  ) |>
  dplyr::filter(flipped_pct_abs)

# after make-signed-bars-data, set eval: true on this once
#| label: export-flip-tables
#| eval: false
out_tbl <- file.path("outputs","tables"); dir.create(out_tbl, recursive = TRUE, showWarnings = FALSE)

mk_pct_abs <- function(df, suffix) {
  df |> tidyr::pivot_wider(names_from = metric, values_from = mean_chg) |>
    dplyr::mutate(flipped_pct_abs = sign(pct) != sign(abs) & !is.na(pct) & !is.na(abs)) |>
    dplyr::filter(flipped_pct_abs) |>
    dplyr::rename(mean_abs = abs, mean_pct = pct) |>
    dplyr::mutate(source = suffix)
}

```

```

}

flips_keep_trim    <- mk_pct_abs(signed_keep0,           "keep_trim")
flips_drop_trim    <- mk_pct_abs(signed_drop0,          "drop_trim")
flips_keep_notrim  <- mk_pct_abs(signed_keep0_no_trim, "keep_notrim")
flips_drop_notrim  <- mk_pct_abs(signed_drop0_no_trim, "drop_notrim")

# save
saveRDS(list(flips_keep_trim, flips_drop_trim, flips_keep_notrim, flips_drop_notrim),
        file.path(out_tbl, "flips_pct_abs_all.rds"))
readr::write_csv(dplyr::bind_rows(flips_keep_trim, flips_drop_trim,
                                  flips_keep_notrim, flips_drop_notrim),
                 file.path(out_tbl, "flips_pct_abs_all.csv"))

```

6 Load cached long table (`plt_long`) for hotspots/KS

Hotspot export and KS still expect `plt_long`. To avoid the heavy pivot here, we load the cached table from `outputs/tables/plt_long.rds`. If it's missing, run `scratch/pivot_long.qmd` (or `scratch/pivot_long.R`) to regenerate it.

7 Hotspot extraction workflow (global + subregional)

We identify per-service hotspots using a 5% percentile rule and **direction vectors**: - `loss_services` (e.g., Nature_Access, Pollination, N/Sed_Ret_Ratio, C_Risk_Red_Ratio) → we flag the **lowest** values; - `gain_services` (Sed_export, N_export, C_Risk) → we flag the **highest** values.

We run this **once globally** and then **once per subregion** (World Bank region, income group, continent, UN region, WWF biome). For each run and for each metric (absolute and percent change) we write a compact **GPKG** containing only the hotspot cells, plus a CSV index:

- Output root: `processed/hotspots/`
 - `abs/global/hotspots_global_abs.gpkg`
 - `pct/global/hotspots_global_pct.gpkg`
 - `abs/<group_col>/hotspots_<group_col>_<group_val>_abs.gpkg`
 - `pct/<group_col>/hotspots_<group_col>_<group_val>_pct.gpkg`
- Index: `processed/hotspots/_hotspots_index.csv` (columns: scope, group_col, group_val, metric, n_hot, gpkg).

These files are meant for QGIS/QA and downstream stats (e.g., KS) without recomputing hotspots.

7.1 Hotspot rules & export configuration

The analysis uses a single, central configuration so the hotspot rules are consistent everywhere:

Thresholding: we flag hotspots using the top/bottom tails of the distribution per service. Here we use a percentile cutoff (e.g., 5%) rather than a fixed count.

Direction of concern: services in loss are “worse when they go down” (we keep the lowest tail); services in gain are “worse when they go up” (we keep the highest tail).

Combos (optional): grouped service sets that we count per cell for quick composite summaries.

Export switches: choose whether to write GPKGs and/or the CSV index.

```
HOTS_CFG <- list(
  analysis_name      = "global_NCP_hotspots",
  pct_cutoff         = 0.05,
  threshold_mode    = "percent",
  rule_mode          = "vectors",
  loss               = c("Nature_Access", "Pollination", "N_Ret_Ratio", "Sed_Ret_Ratio", "C_Risk_Red_Ratio"),
  gain               = c("Sed_export", "N_export", "C_Risk"),
  combos             = list(
    deg_combo        = c("Nature_Access", "Pollination", "N_export", "Sed_export", "C_Risk"),
    rec_combo        = c("Nature_Access", "Pollination", "N_Ret_Ratio", "Sed_Ret_Ratio", "C_Risk_Red_Ratio")
  ),
  # centralize the grouping columns here
  groupings          = c("income_grp", "region_wb", "continent", "region_un", "WWF_biome"),
  # IO
  write_layers       = TRUE,
  write_index        = TRUE,
  out_dir            = file.path(data_dir(), "processed", "hotspots")
)
```

::: callout-note

Hotspot configuration

- Cutoff: 5% (percent)
 - Rule mode: vectors
 - Loss services: Nature_Access, Pollination, N_Ret_Ratio, Sed_Ret_Ratio, C_Risk_Red_Ratio
 - Gain services: Sed_export, N_export, C_Risk
 - Combos: **deg_combo**: Nature_Access, Pollination, N_export, Sed_export, C_Risk
**rec_com...
 - Groupings: income_grp, region_wb, continent, region_un, WWF_biome
 - ...- Write layers: TRUE | Write index: TRUE
 - Output dir: `/home/jeronimo/data/global_ncp/processed/hotspots`
- :::

7.2 Validate hotspot configuration

```
# stopifnot(exists("plt_long"))
# svc_all <- unique(plt_long$service)
# if (length(intersect(HOTS_CFG$loss, HOTS_CFG$gain)) > 0) {
#   stop("A service appears in BOTH `loss` and `gain`. Fix HOTS_CFG.")
# }
# miss_loss <- setdiff(HOTS_CFG$loss, svc_all)
# miss_gain <- setdiff(HOTS_CFG$gain, svc_all)
# if (length(miss_loss) > 0 || length(miss_gain) > 0) {
#   warning("Services in HOTS_CFG not found in `plt_long$service`:\n",
#         if (length(miss_loss)) paste0(" - missing loss: ", paste(miss_loss, collapse=", ")),
#         if (length(miss_gain)) paste0(" - missing gain: ", paste(miss_gain, collapse=", ")))
# }
```

7.3 Export hotspot layers

Note

Hotspot export module

- Computes hotspot cells once (global + by subregion) for ABS and PCT change.
- Writes compact GPKGs for mapping/QA and maintains `_hotspots_index.csv`.
- Prereqs: `plt_long` in memory, `HOTS_CFG` defined (loss/gain/combos/etc.).

8 Trimmed change bar plots

Note

How to read these bars

- Each bar shows the **trimmed mean absolute change** ($|\Delta|$) per service within each group; facet axes are free.
- Bars are **always positive** by design: height = **magnitude of change**, not direction.
- Direction-of-concern used elsewhere in the analysis:
 - Worse when **up** `Sed_export`, `N_export`, `C_Risk`.
 - Worse when **down** `Nature_Access`, `Pollination`, `N_Ret_Ratio`, `Sed_Ret_Ratio`, `C_Risk_Red_Ratio`.
- These bars answer “**where is change largest?**”. See hotspot maps/violins for **up** vs. **down** patterns.

Short answer: your current barplots use all grid cells (the full 10-km population), not just hotspots. They summarize trimmed means per subregion/global from `plt_long` via `aggregate_change_simple()`,

with `cut_q=0.999` to cap outliers and (optionally) `drop_zeros=TRUE`. That's why every bar is positive—those bars are the magnitude of change, not the direction.

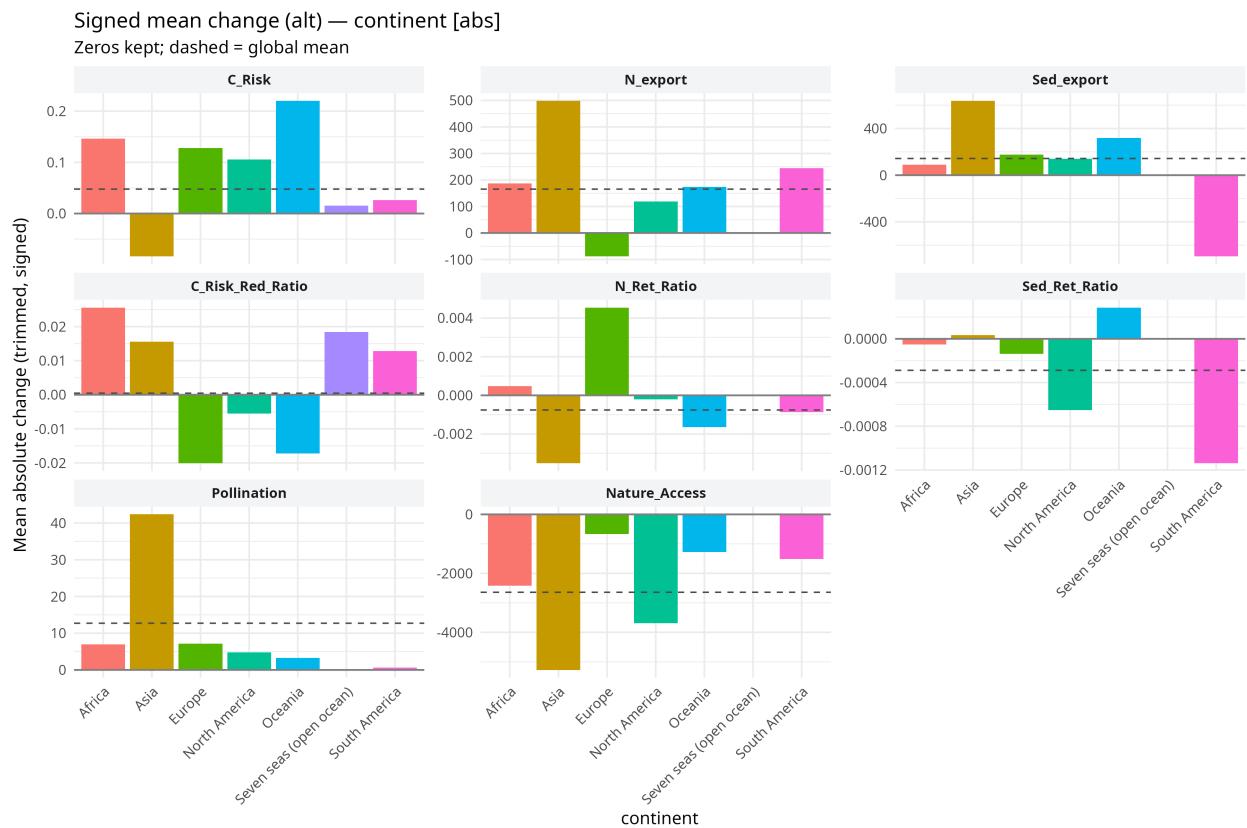
Outputs: PNGs land in `outputs/plots/{abs|pct}/<group_col>/bars_*.png` plus a flat copy in `outputs/plots/latest/bars/` for embedding here. The legacy global-only single-bar chart was removed to keep the gallery focused on grouped views.

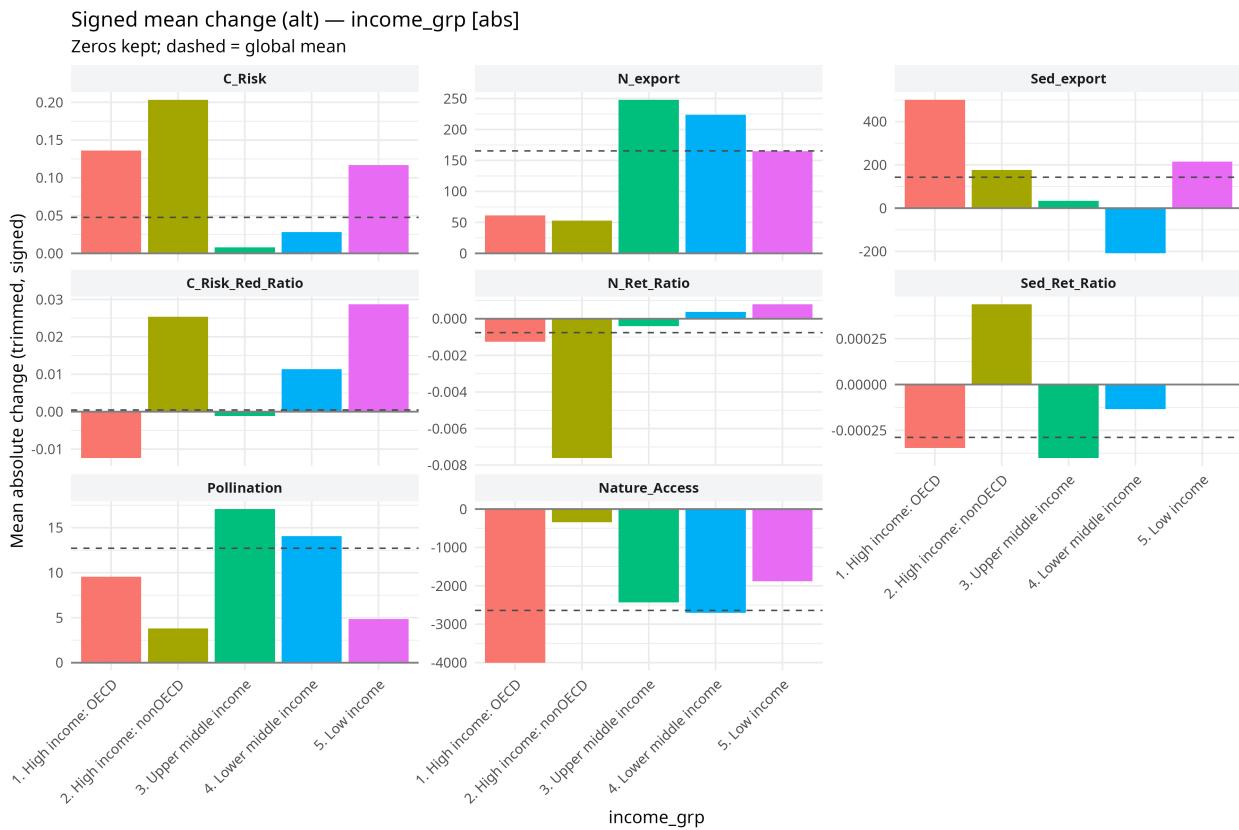
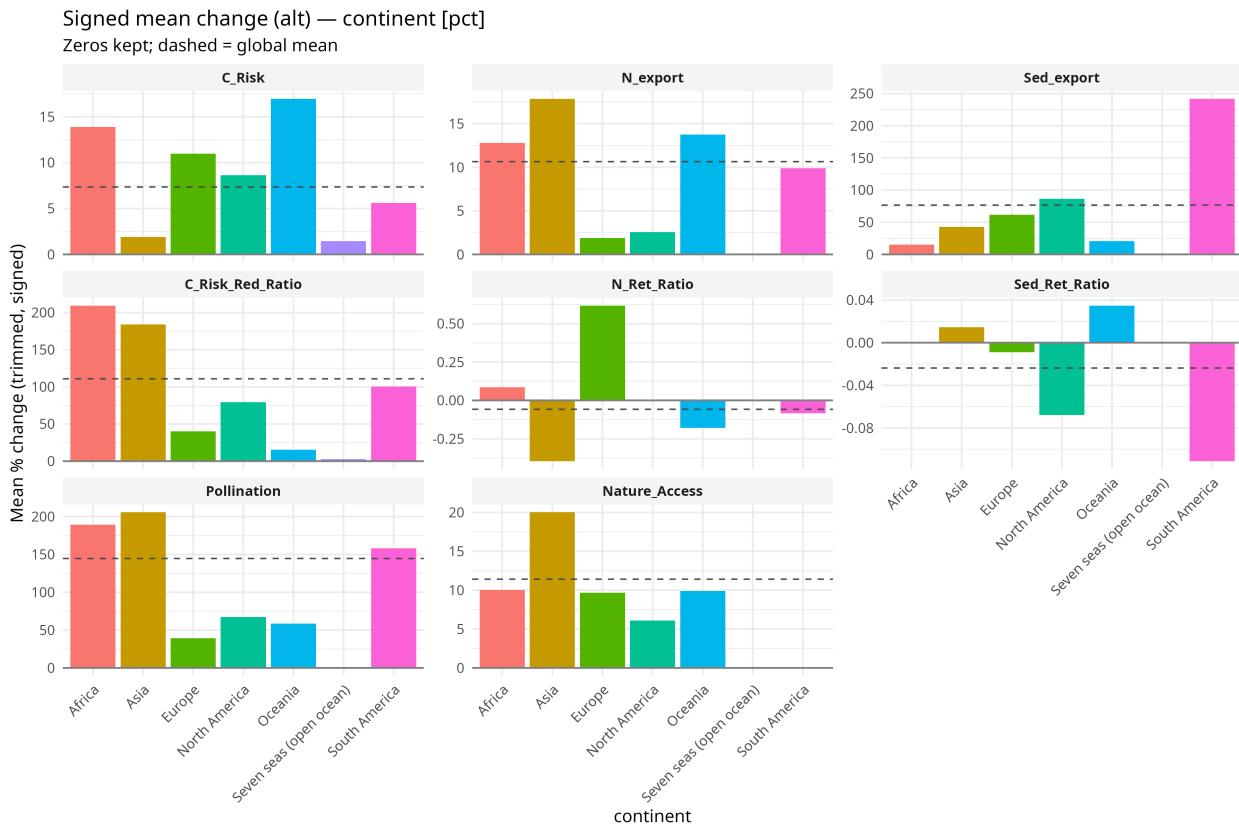
Generate trimmed-mean bar plots (abs & pct) for each grouping. Plots save to `outputs/plots/...` and are embedded below.

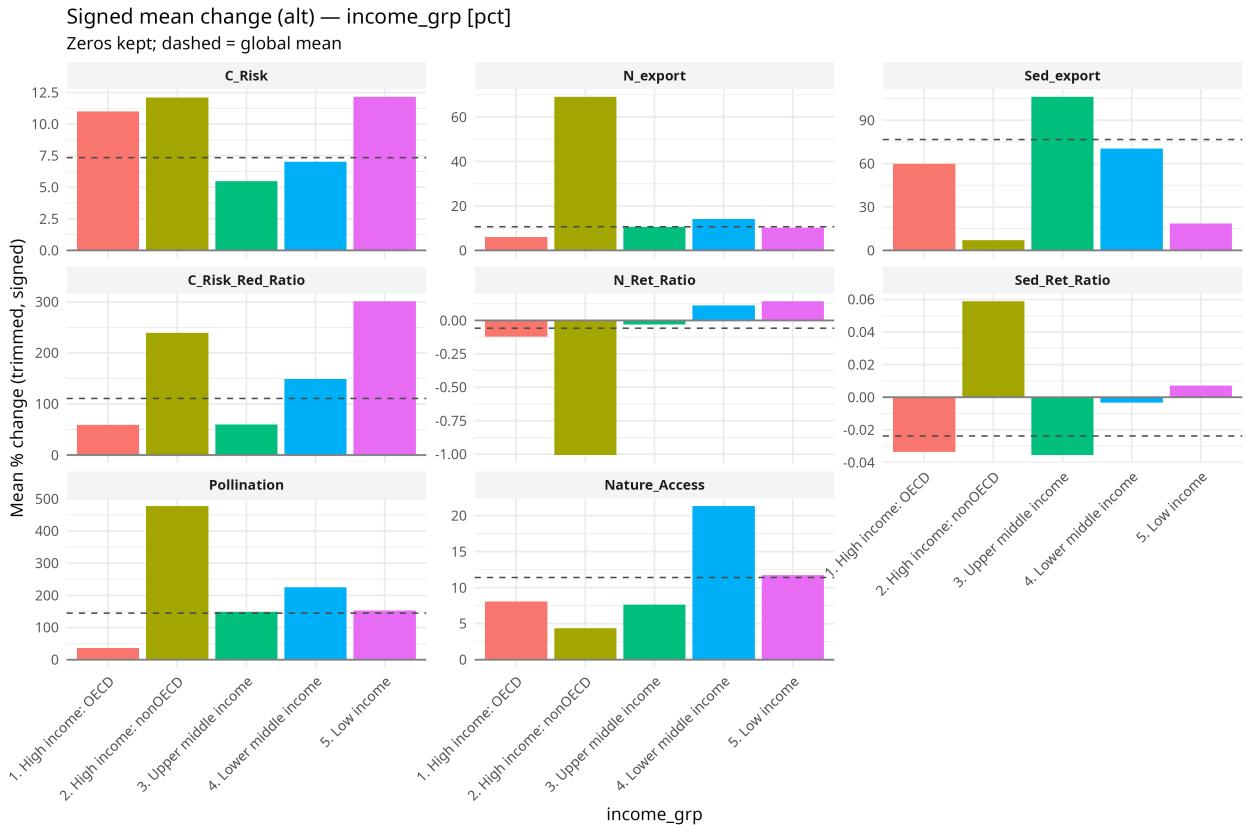
8.1 TODO: Hotspot vs non-hotspot contribution bars

Future idea: stack bars to show how much of each group's total signed change comes from hotspot cells vs the complement. Steps: compute group means separately for hotspot and non-hotspot subsets (using the same 5% rules), then plot stacked bars (hotspot share + non-hotspot share) per service/group. Keep both absolute and percent-change variants; avoid overwriting existing plots (write to `outputs/plots/signed_hotshare/`).

8.2 Signed means (alt variants: `keep0` vs `drop0`; dashed = global mean)

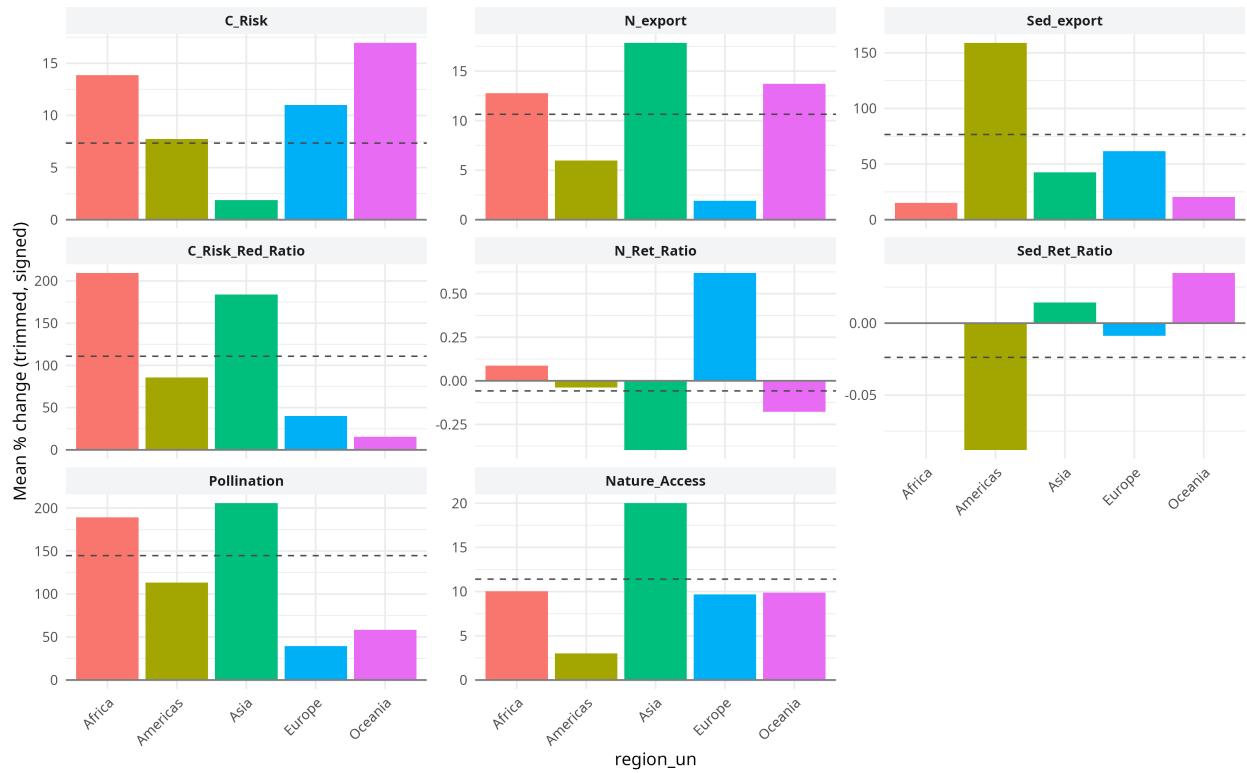






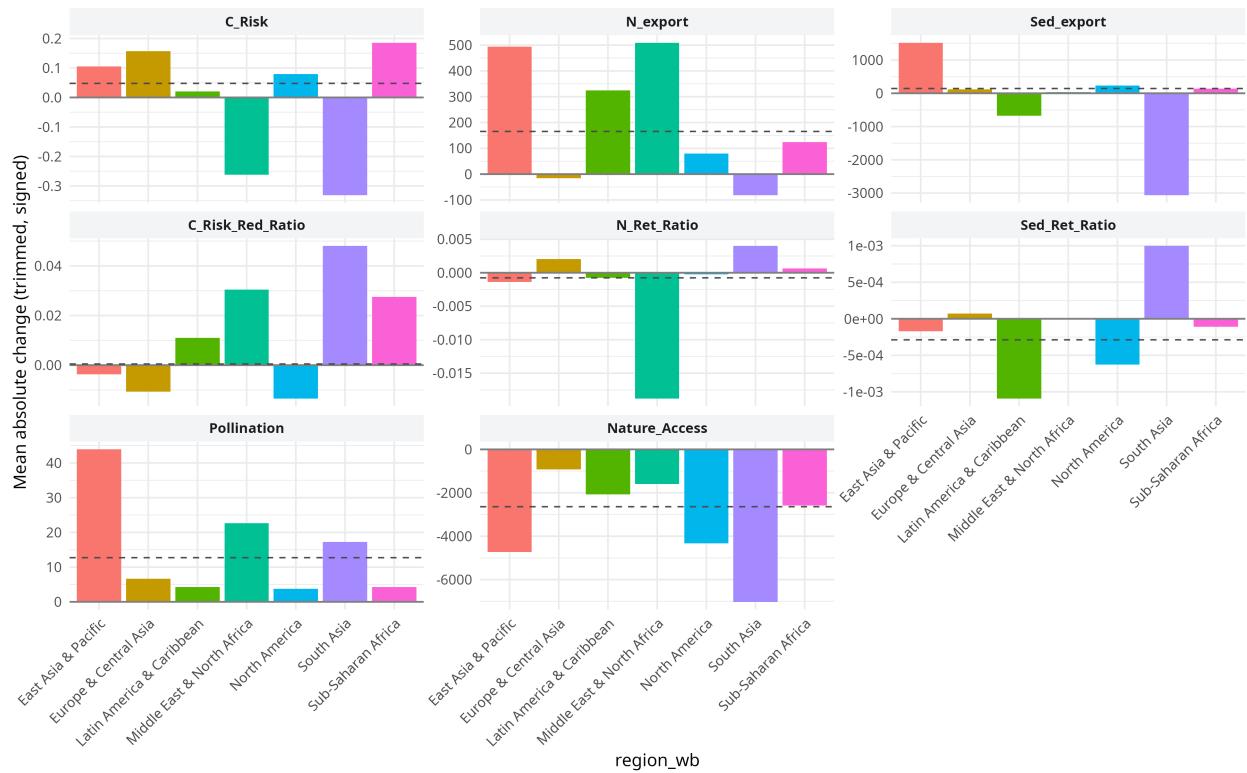
Signed mean change (alt) — region_un [pct]

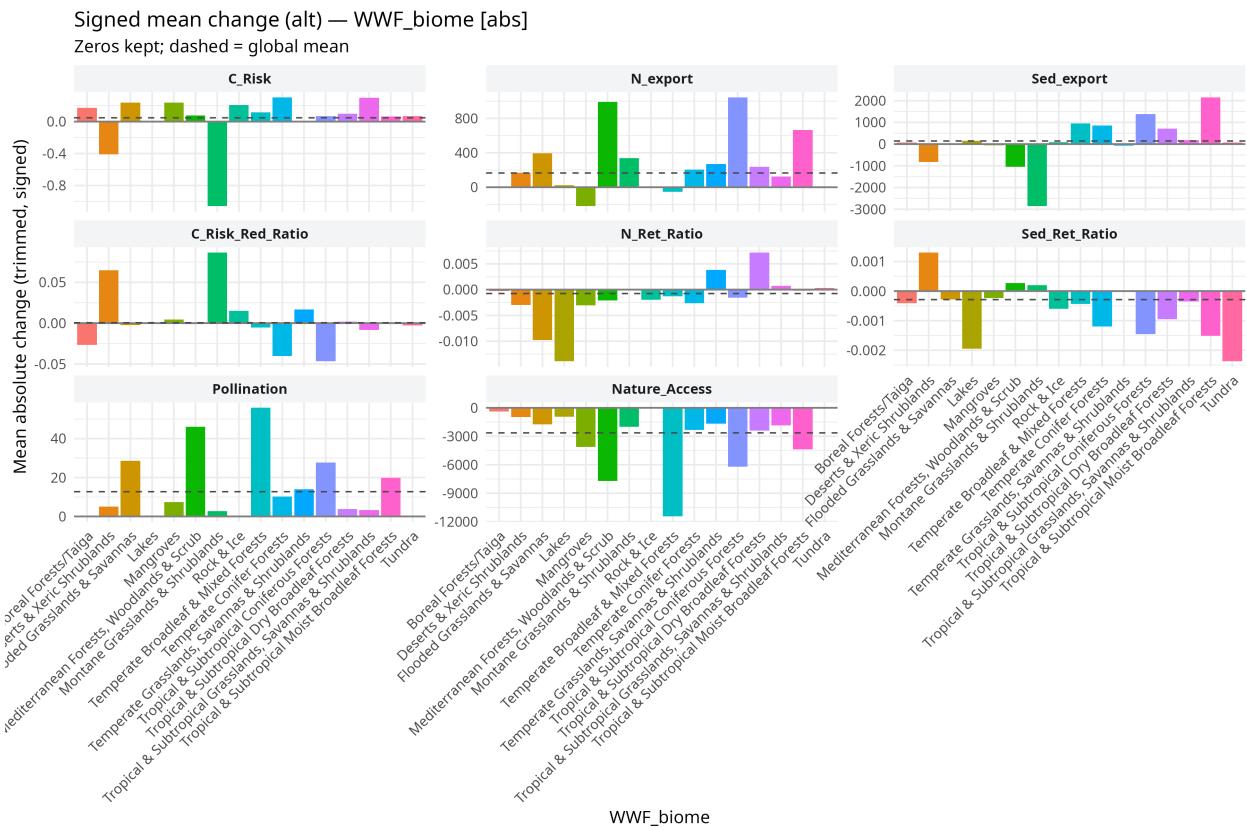
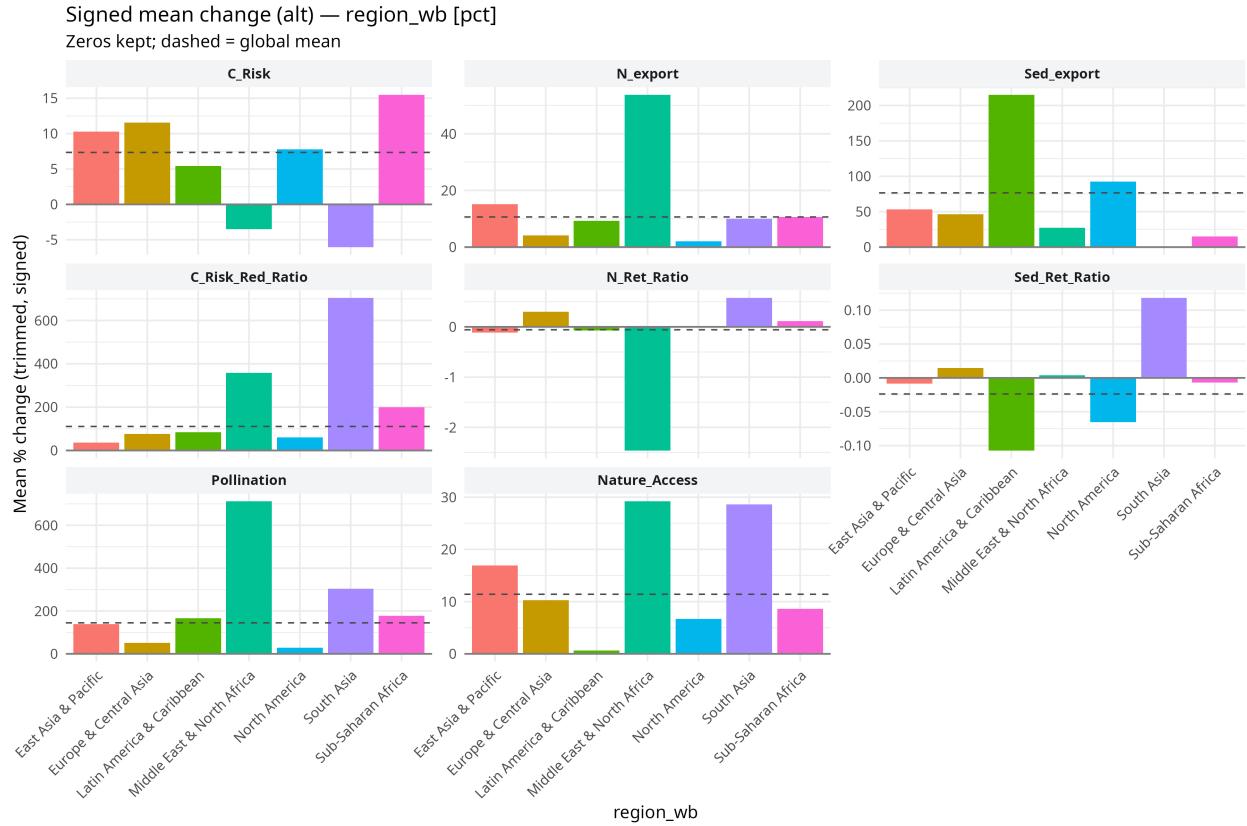
Zeros kept; dashed = global mean

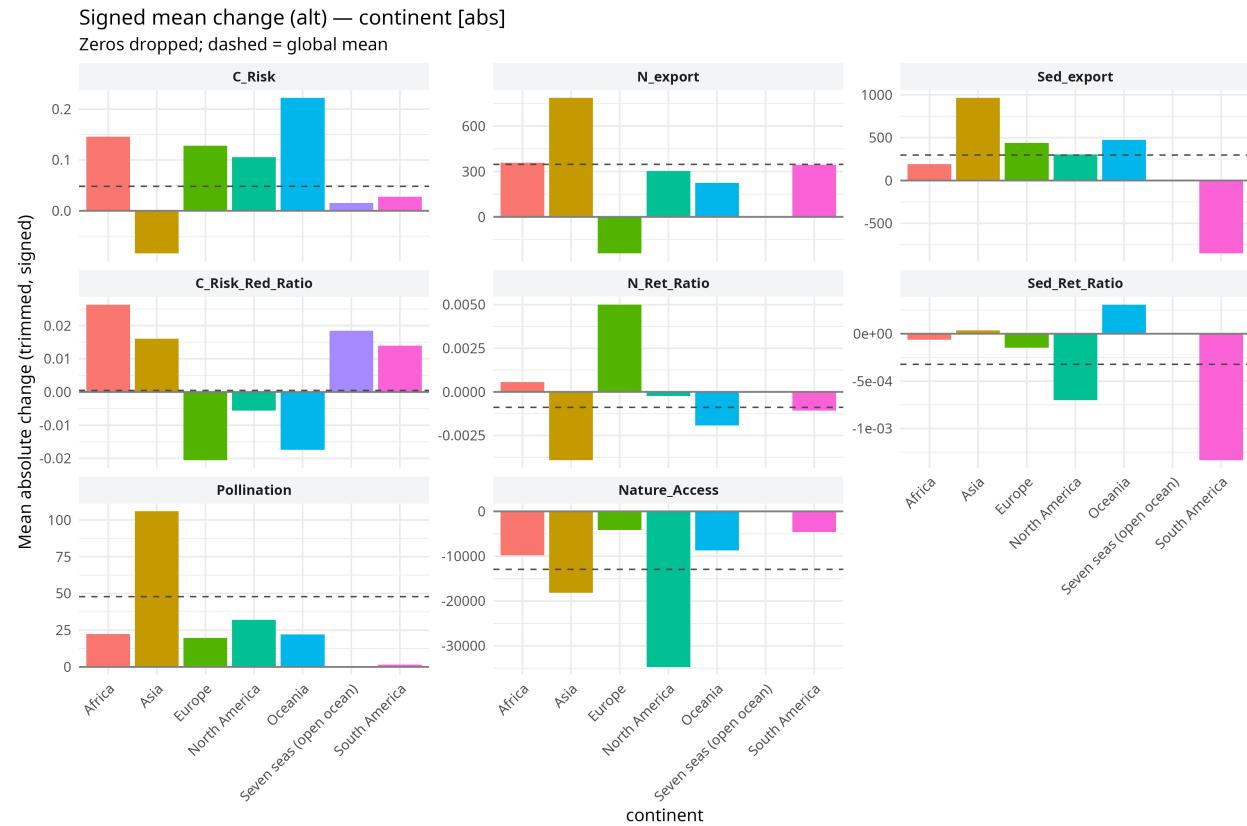


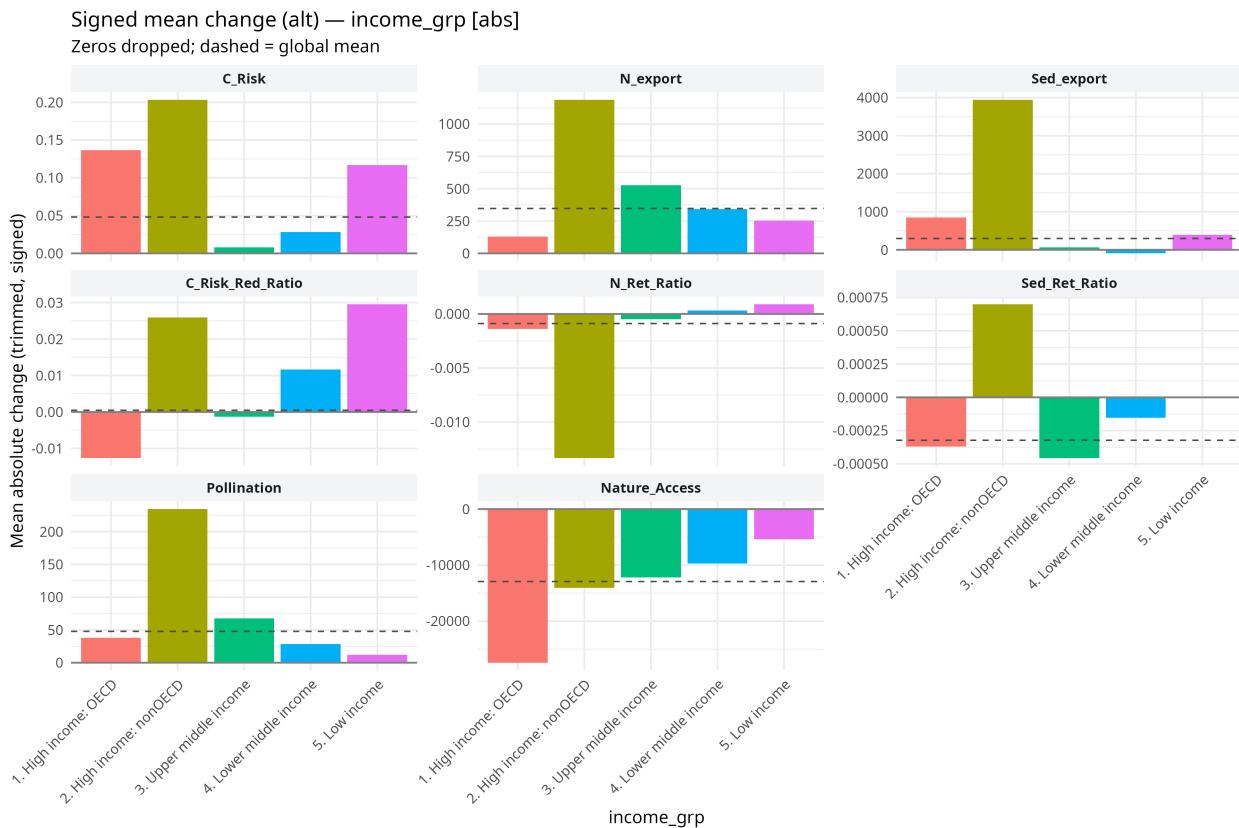
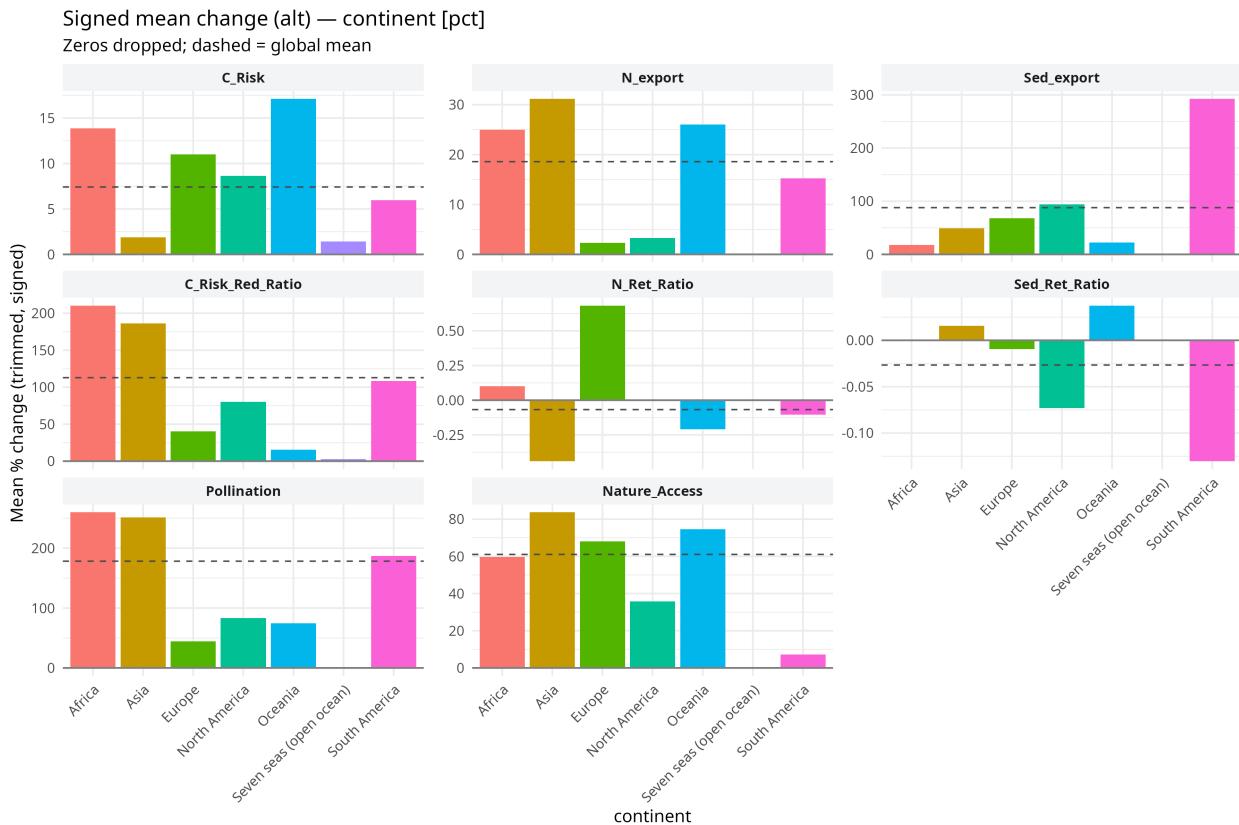
Signed mean change (alt) — region_wb [abs]

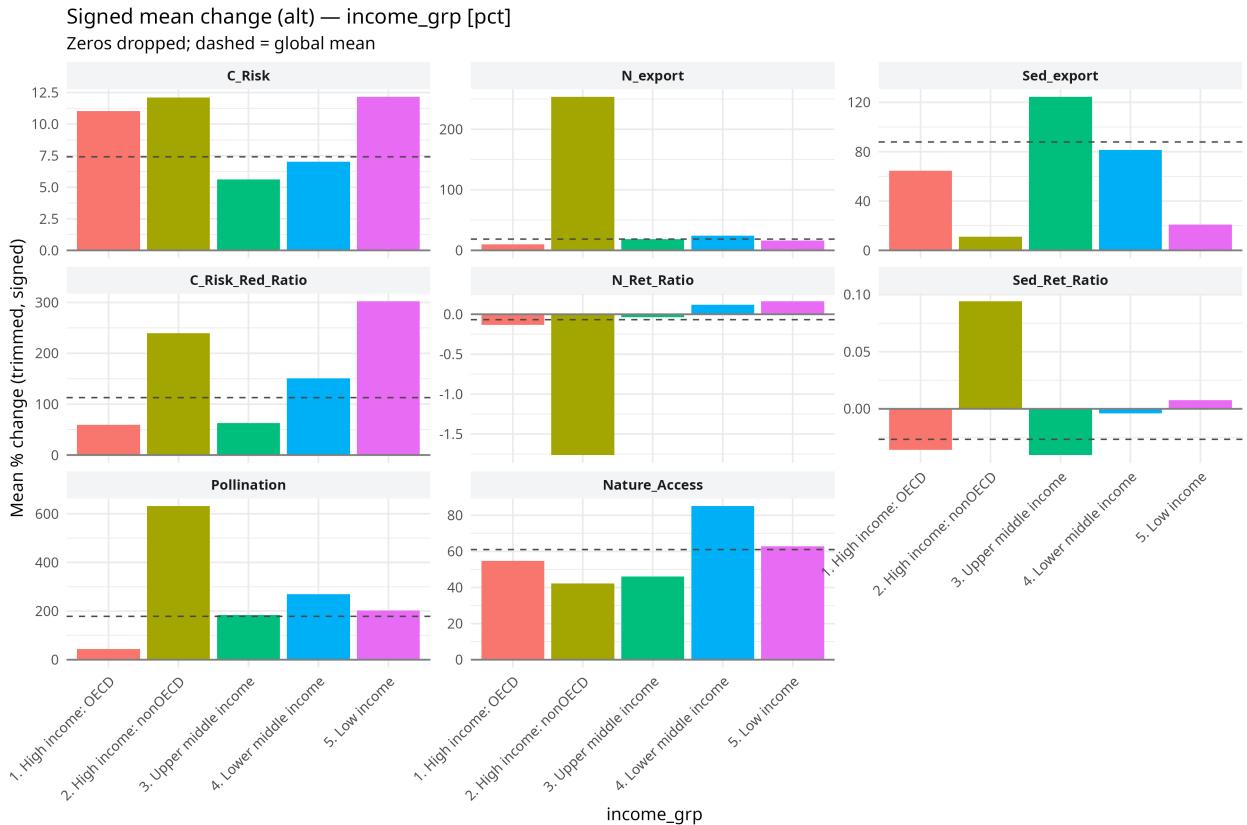
Zeros kept; dashed = global mean





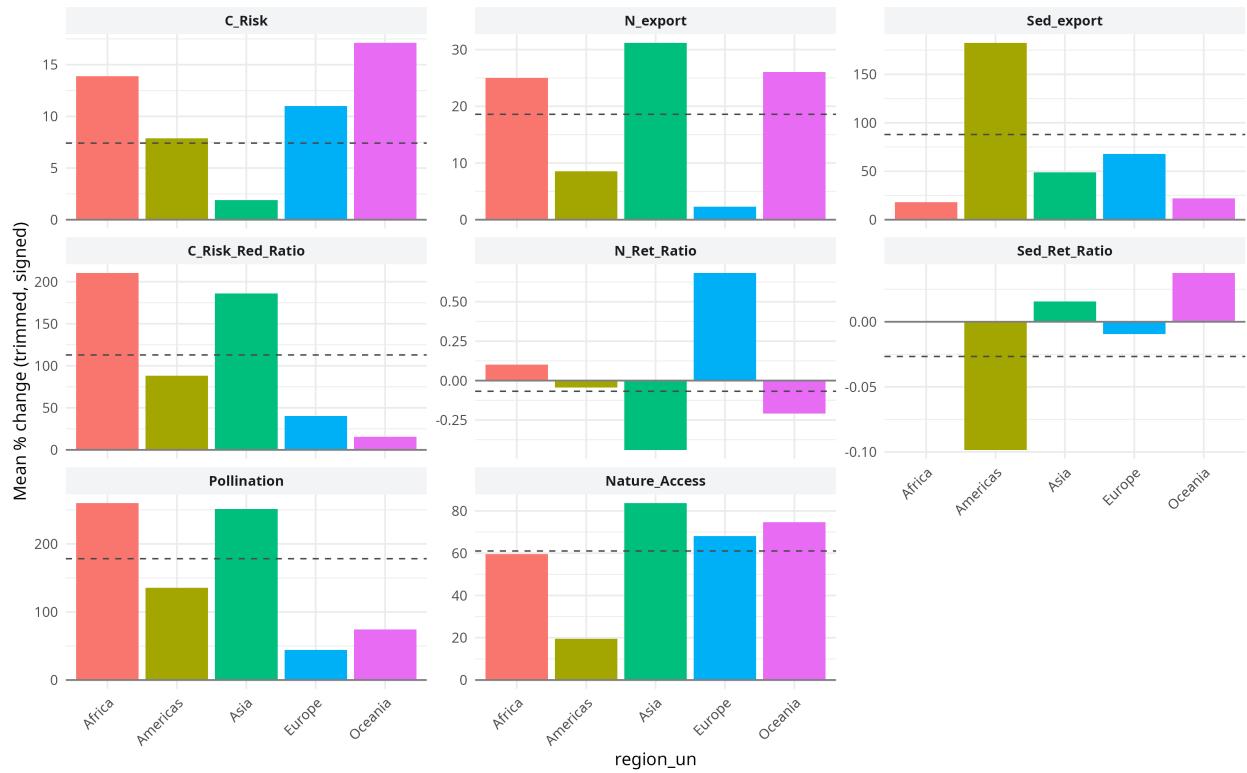






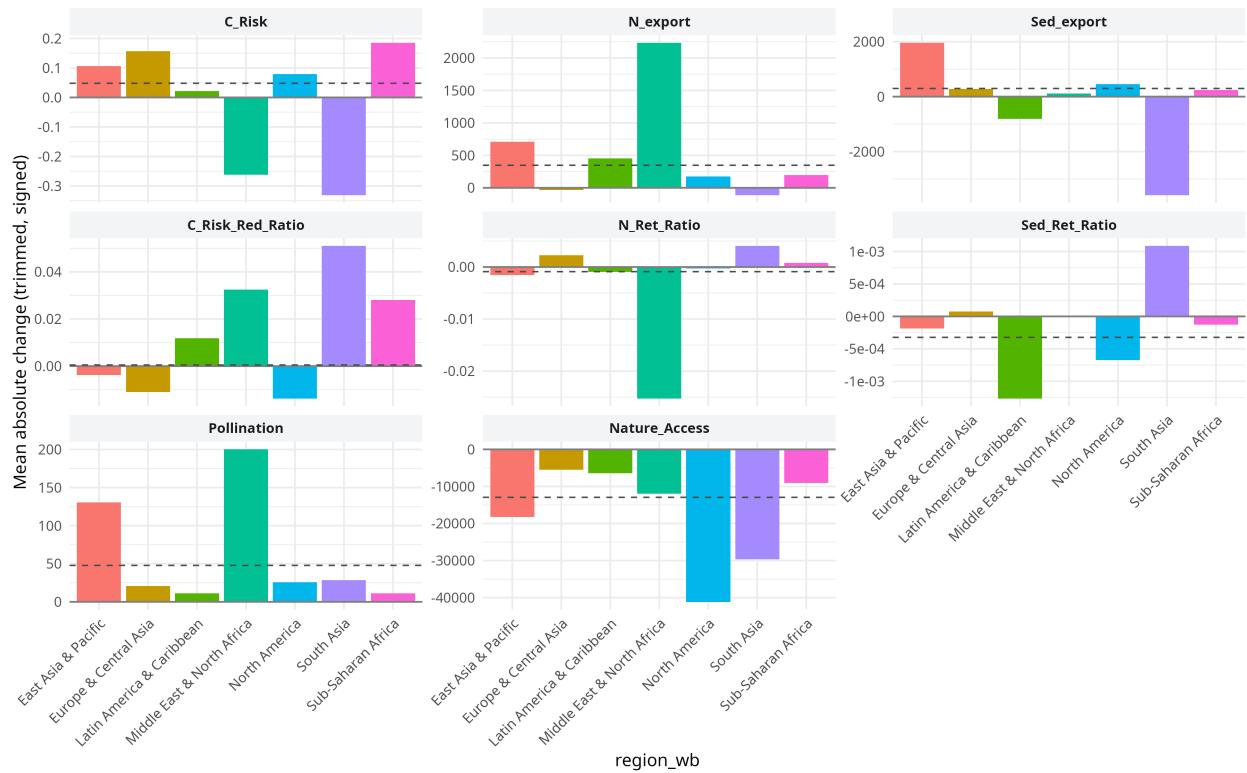
Signed mean change (alt) — region_un [pct]

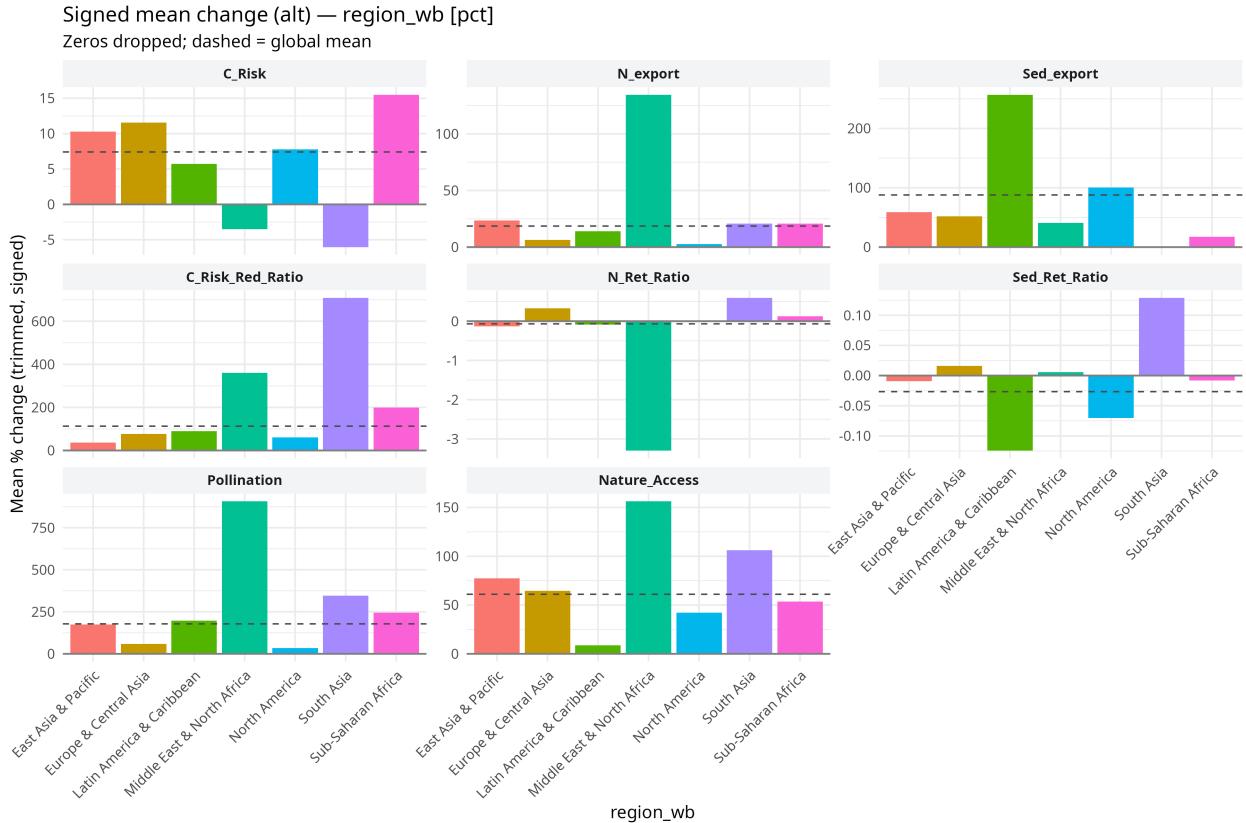
Zeros dropped; dashed = global mean

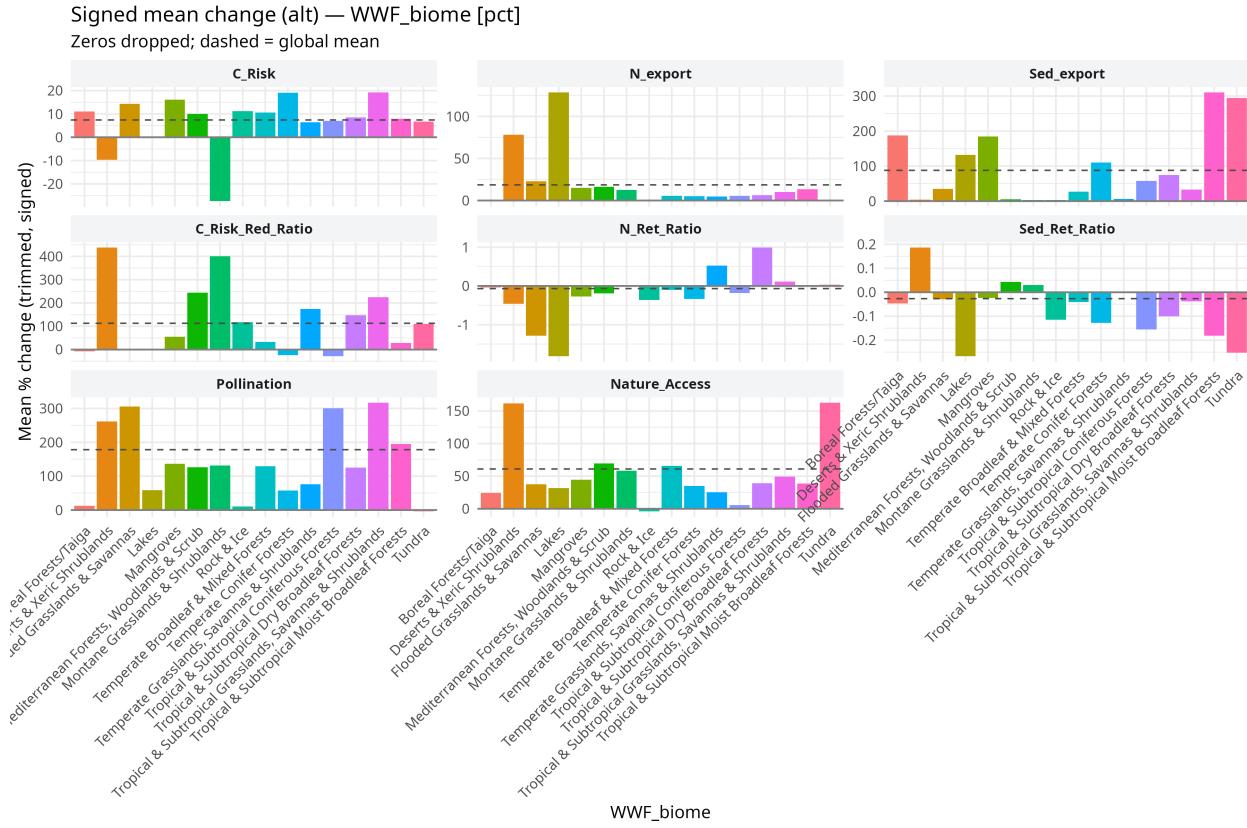


Signed mean change (alt) — region_wb [abs]

Zeros dropped; dashed = global mean





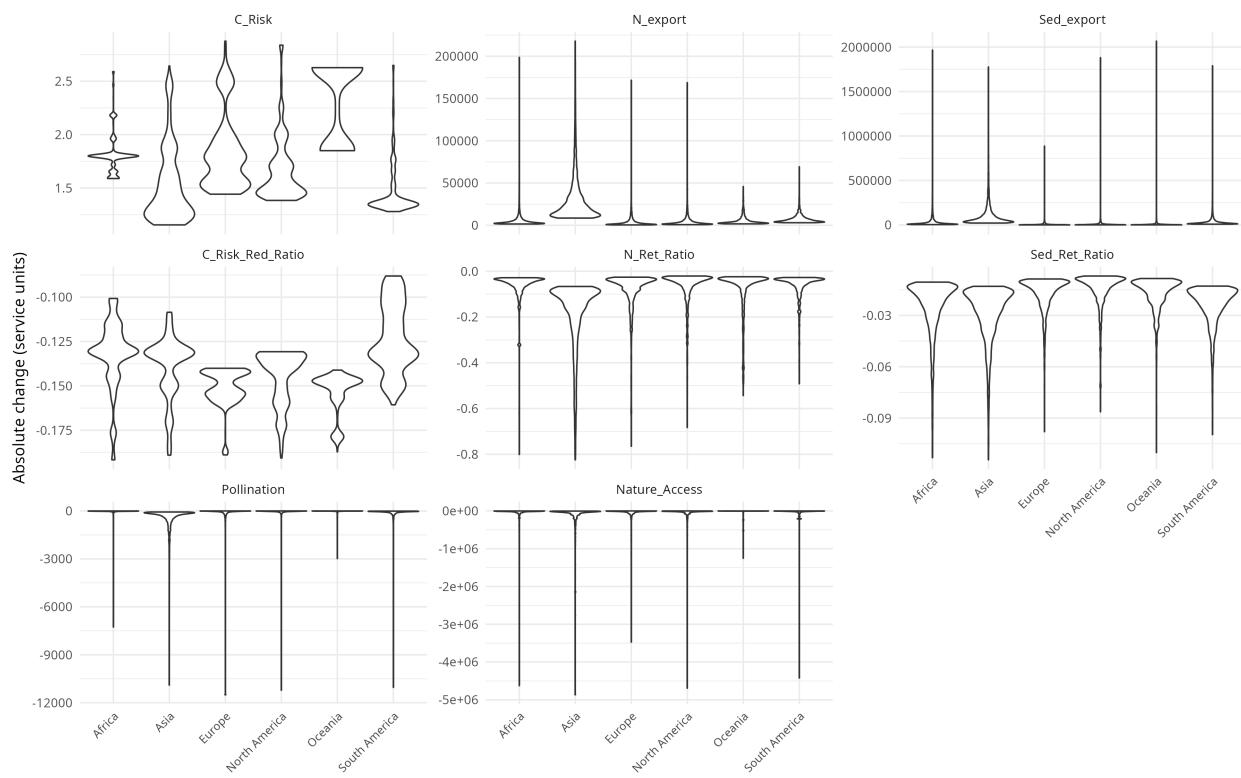


9 Hotspot violin plots

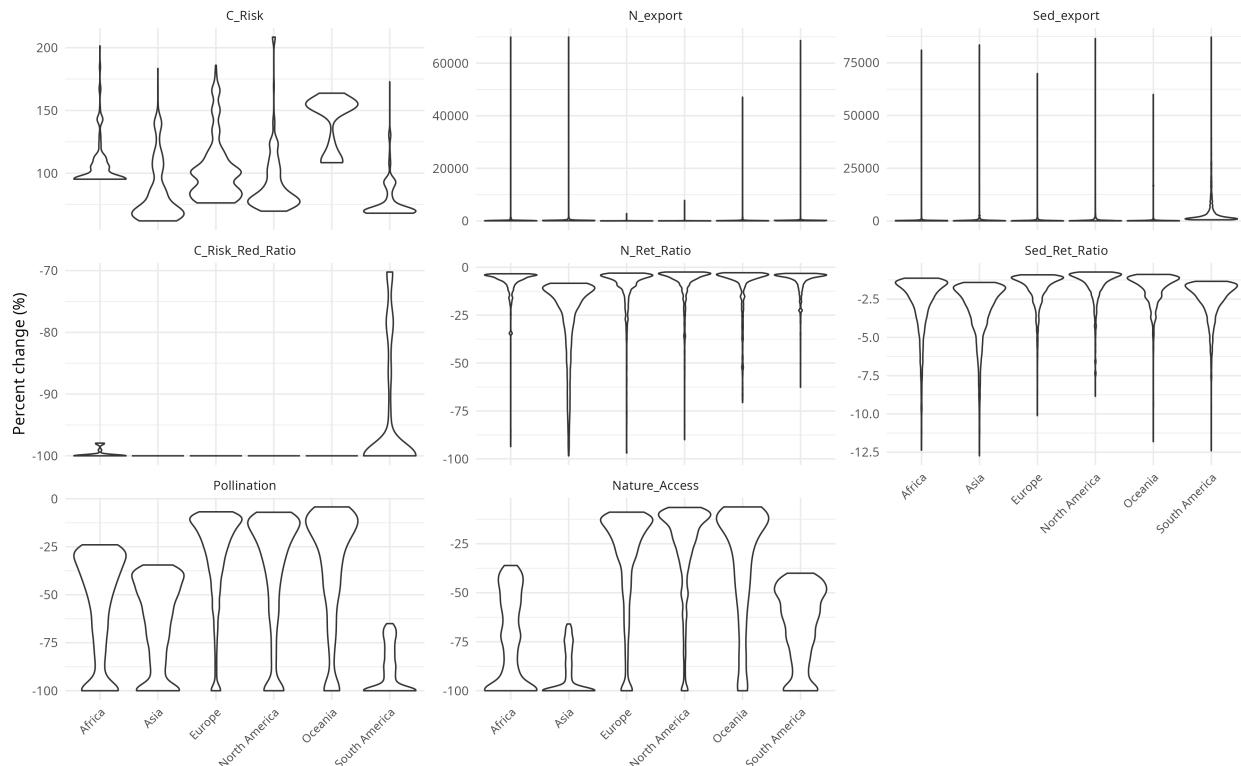
Summarize hotspot distributions by group using saved PNGs; skip computation if group columns are absent. Helper `run_hotspot_violins_by()` now lives in `R/hotspot_violins.R`, so we just call it from this report.

Outputs: PNGs are written to `outputs/plots/{abs|pct}/<group_col>/violins_*.png` and mirrored into `outputs/plots/latest/violins/` for embedding.

Absolute change in hotspots by continent
Violins only · NA/0 removed · per-service 99.9% trim

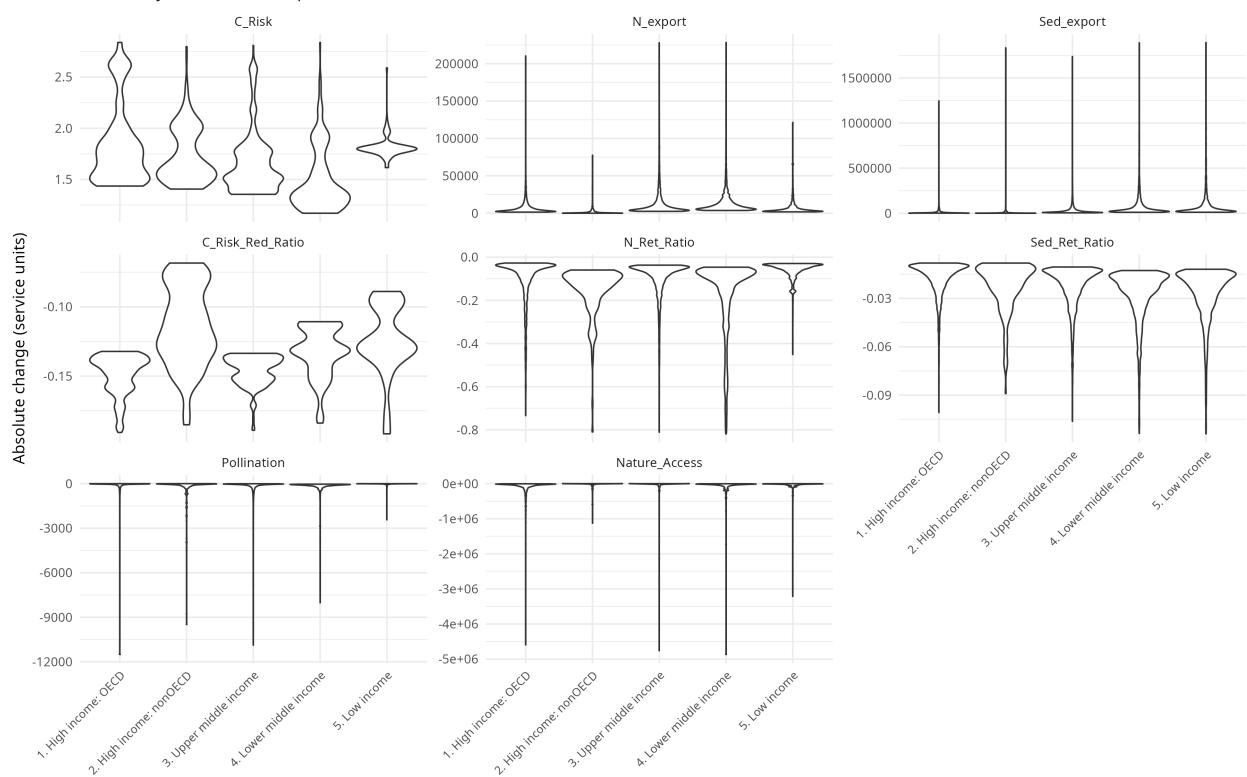


Percent change in hotspots by continent
Violins only · NA/0 removed · per-service 99.9% trim



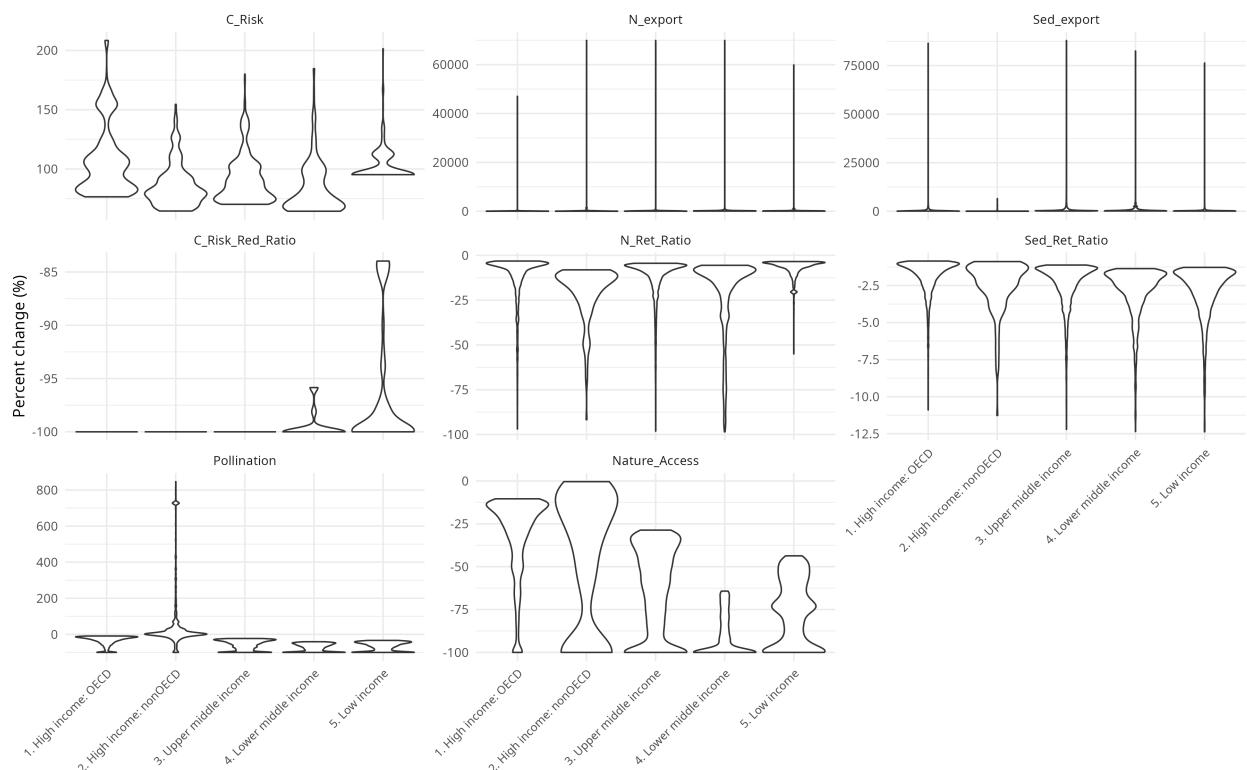
Absolute change in hotspots by income_grp

Violins only · NA/0 removed · per-service 99.9% trim

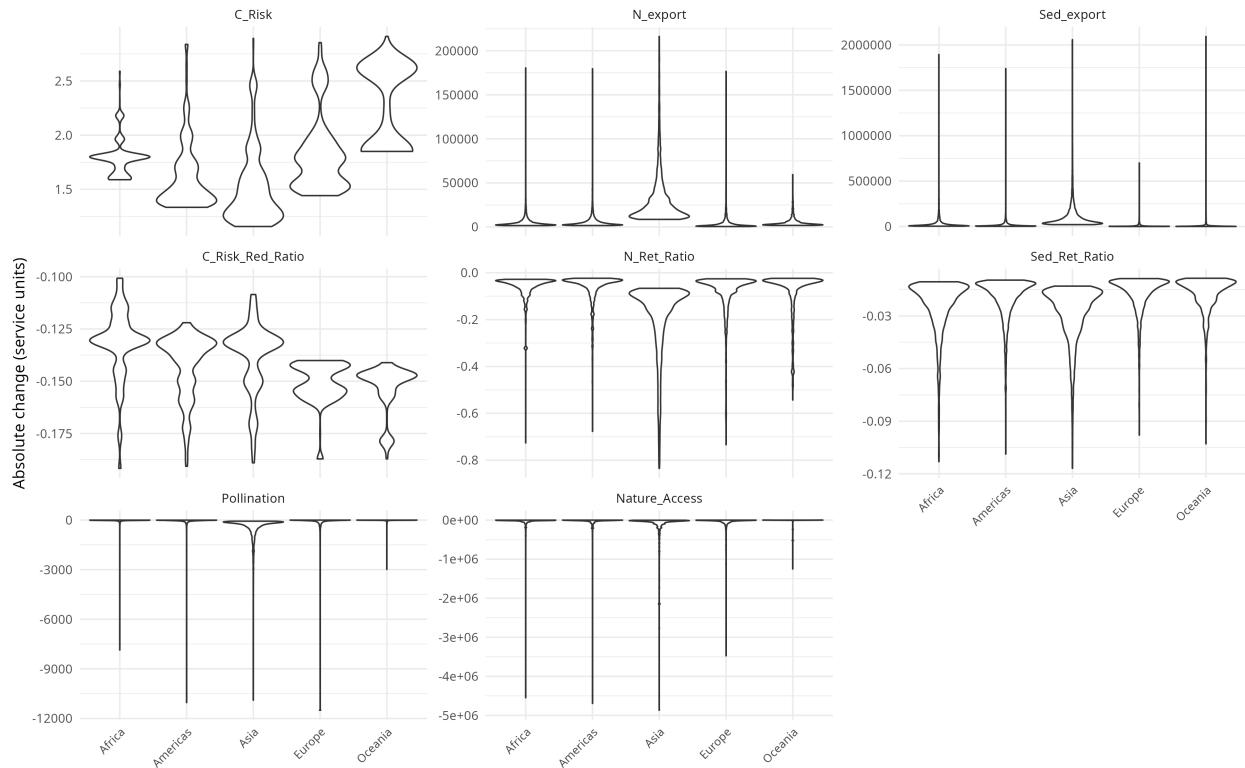


Percent change in hotspots by income_grp

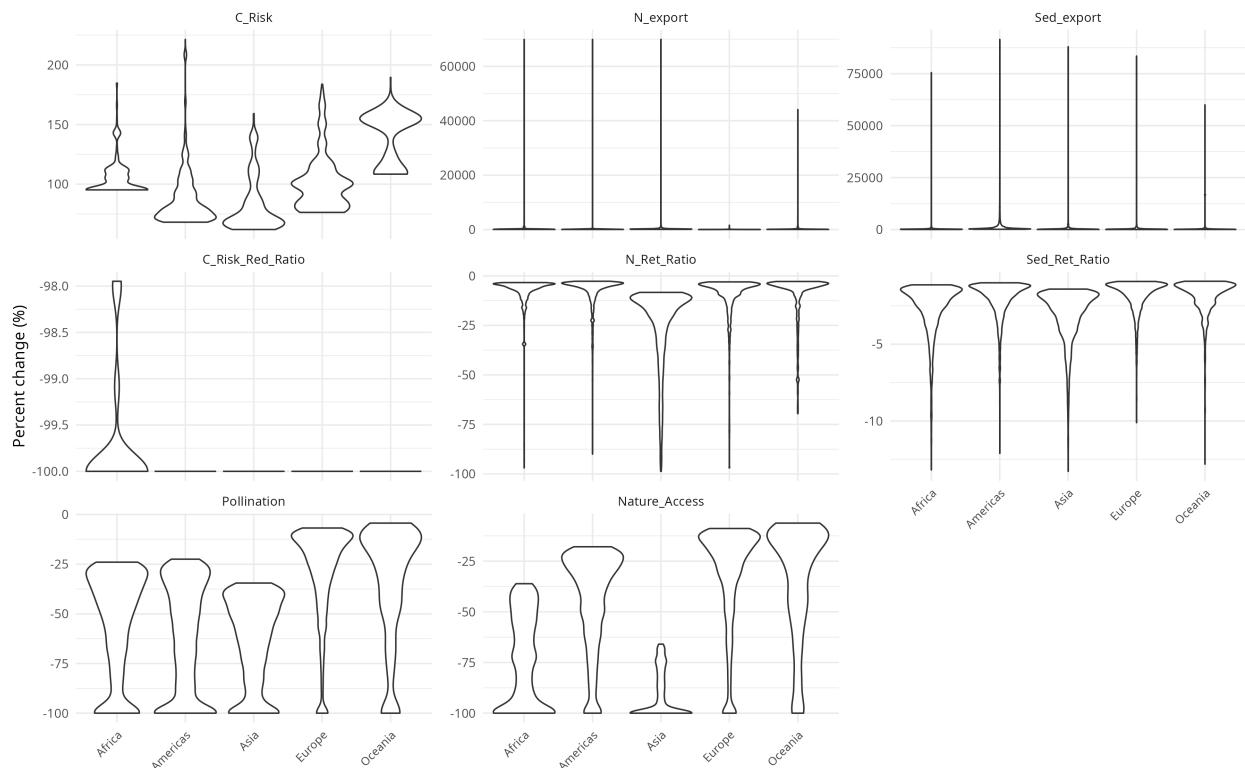
Violins only · NA/0 removed · per-service 99.9% trim



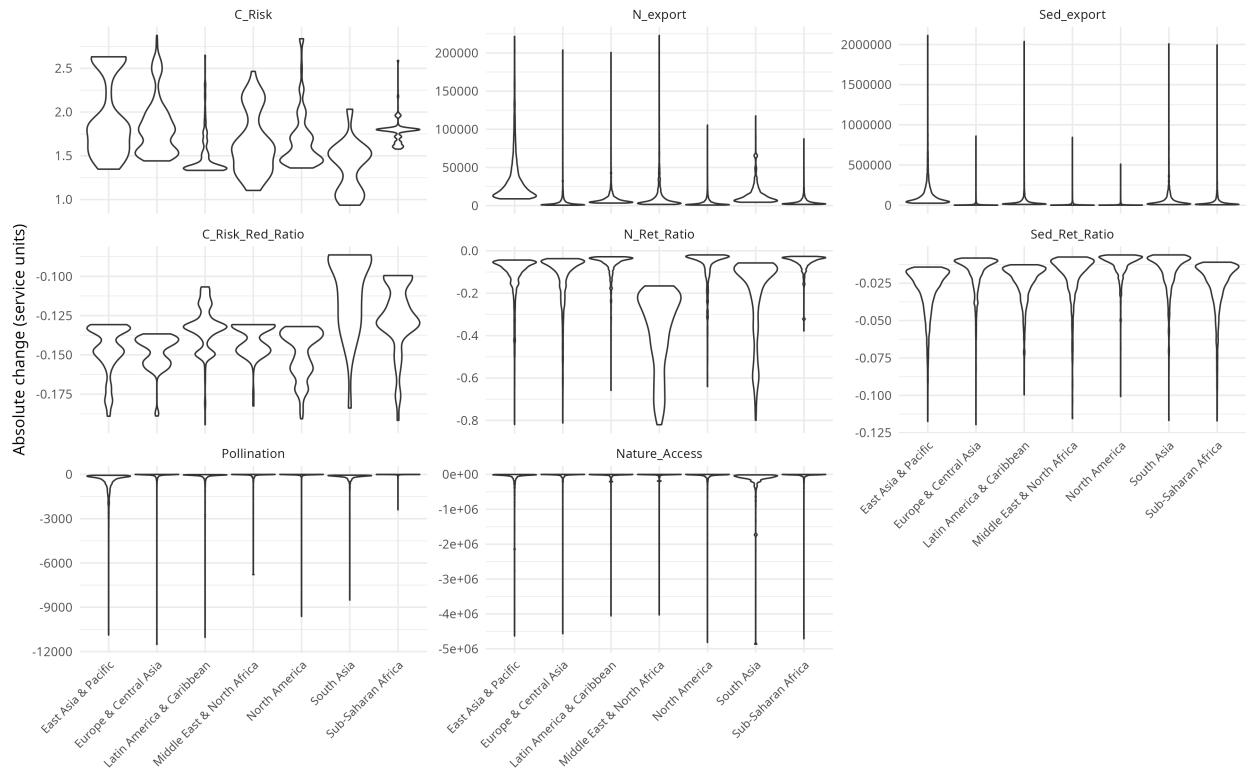
Absolute change in hotspots by region_un
Violins only · NA/0 removed · per-service 99.9% trim



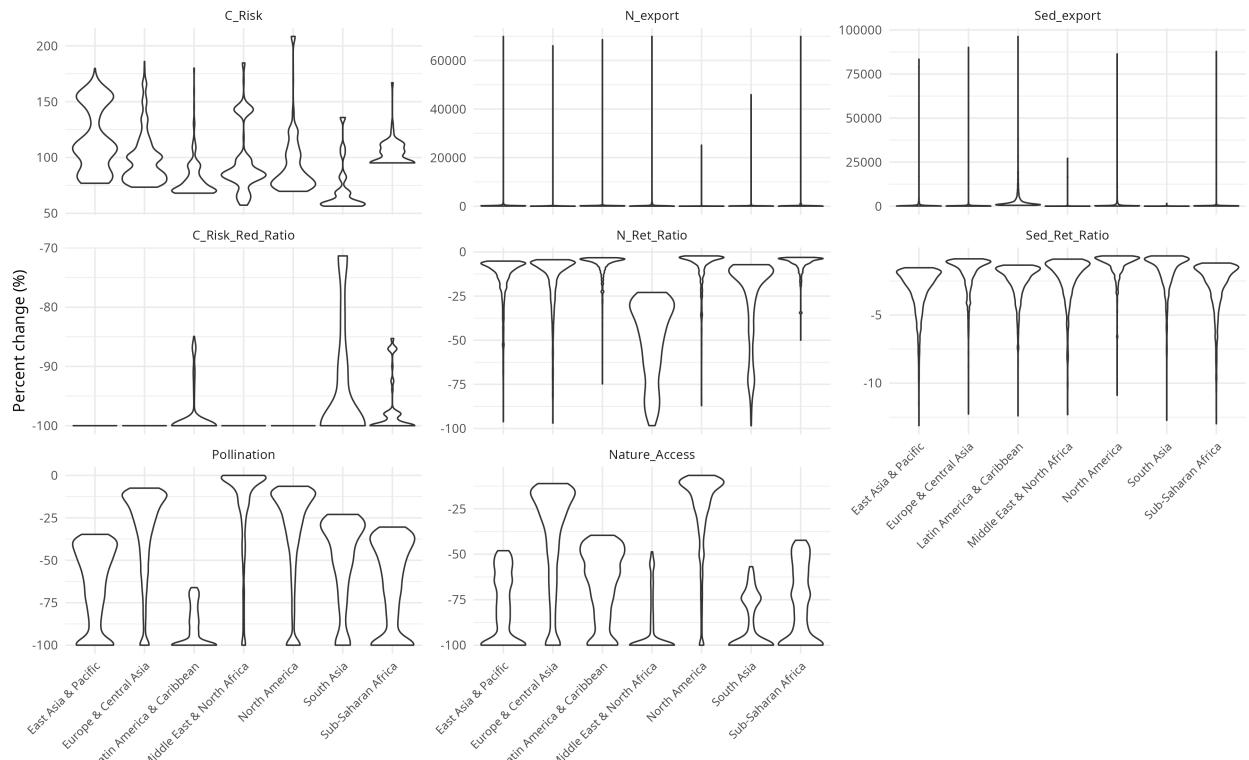
Percent change in hotspots by region_un
Violins only · NA/0 removed · per-service 99.9% trim



Absolute change in hotspots by region_wb
Violins only · NA/0 removed · per-service 99.9% trim

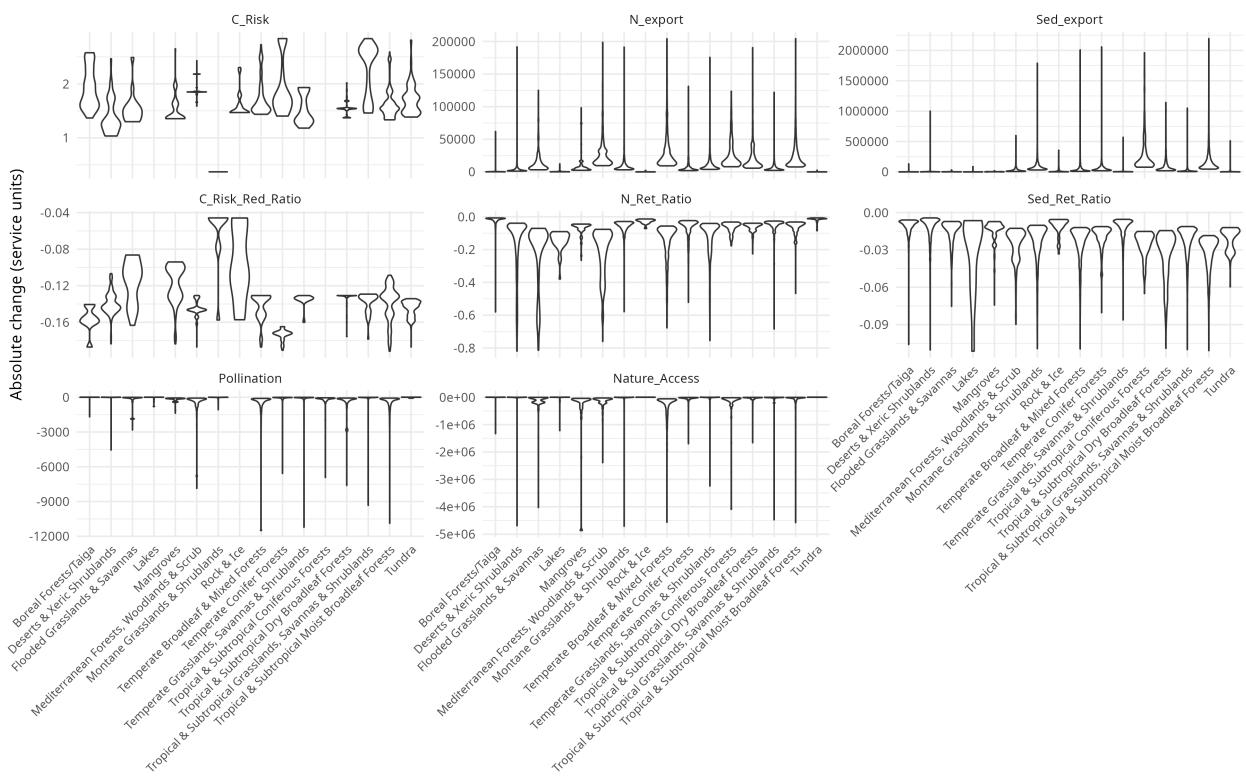


Percent change in hotspots by region_wb
Violins only · NA/0 removed · per-service 99.9% trim



Absolute change in hotspots by WWF biome

Violins only · NA/0 removed · per-service 99.9% trim



Percent change in hotspots by WWF biome

Violins only · NA/0 removed · per-service 99.9% trim

